

UNIVERSITY COLLEGE CORK

DEPARTMENT OF COMPUTER SCIENCE  
FINAL YEAR PROJECT BSc IN COMPUTER  
SCIENCE

# Word of the People

*Conor Tak Saruwatari*

supervised by  
Dr. James DOHERTY

April, 2017

# DECLARATION OF ORIGINALITY

In signing this declaration, you are conforming, in writing, that the submitted work is entirely your own original work, except where clearly attributed otherwise, and that it has not been submitted partly or wholly for any other educational award.

I hereby declare that:

- this is all my own work, unless clearly indicated otherwise, with full and proper accreditation;
- with respect to my own work: none of it has been submitted at any educational institution contributing in any way to an educational award;
- with respect to anothers work: all text, diagrams, code, or ideas, whether verbatim, paraphrased or otherwise modified or adapted, have been duly attributed to the source in a scholarly manner, whether from books, papers, lecture notes or any other students work, whether published or unpublished, electronically or in print.

Signed: . . . . .

Date: . . . . .

## **Abstract**

In the current age Twitter has quickly become one of the most popular forms of Social Media. Its fast pace and the concise nature of its 140 character limit make it an invaluable source of information into the opinions and stance of a massive population. The goal of this project is to create an efficient means of analyzing and compiling this data in an effective manner such that it can be used for market research, recommender systems, predictive AI and any other number of possible applications. This was developed in Python while restricting available resources to those legally usable within corporate development, such as libraries holding an MIT Software License.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Motivation . . . . .	3
1.2	Goals . . . . .	4
<b>2</b>	<b>Analysis</b>	<b>5</b>
2.1	Objectives . . . . .	5
<b>3</b>	<b>Design</b>	<b>6</b>
3.1	Main Interface . . . . .	6
3.2	Text Cleaner . . . . .	6
3.3	Twitter Scraper . . . . .	7
3.4	Twitter Handler . . . . .	7
3.5	Text Analyzer . . . . .	7
3.6	Database . . . . .	7
<b>4</b>	<b>Implementation</b>	<b>8</b>
4.1	Project Planning . . . . .	8
4.2	Disecting Tweet Structure . . . . .	9
4.3	Database Structuting . . . . .	9
4.4	Test Data Collection . . . . .	10
4.5	Text Analysis . . . . .	11
<b>5</b>	<b>Evaluation</b>	<b>12</b>
5.1	Weights . . . . .	12
5.2	Data Collection . . . . .	12
<b>6</b>	<b>Conclusions</b>	<b>14</b>
6.1	Development . . . . .	14

# List of Figures

4.1	First Gantt Chart Draft . . . . .	8
4.2	Final Gantt Chart Draft . . . . .	9
4.3	Finalized Database Schema . . . . .	9
5.1	Graphed Data . . . . .	13

# Chapter 1

## Introduction

Since 2007, Twitter has quickly risen to become one of the most popular forms of Social Media. Its fast pace and 140 character limit provide quick, concise opinions on a wide variety of topics. This makes it an invaluable source of information into the opinions and stances of a massive population.

Development was handled almost entirely in Python 2.7[1] and use of extra libraries was restricted to those usable within industrial or corporate development environments, such as those with an MIT License[2].

### 1.1 Motivation

While working as a Software Engineer at IBM[3] as part of my Work Placement during the summer of 2016 I spent a lot of time working with large amounts of data. Compiling and aggregating large volumes of raw data into usable reports and logical conclusions. I found this work both interesting and engaging, leaving me with a desire to continue into such analytics after my placement concluded. This project has given me the opportunity to explore this subject matter further, deepening and enhancing my own knowledge and experience in data handling and analysis.

Even after my work at IBM, I knew I had only just scratched the surface of big data analytics. I had mapped reports and diagnostics to hard coded conclusions but I had done very little work into analyzing specific elements of the data itself. Feeling both confident in my ability I figured text analysis would be an interesting topic to explore. However, knowing very little of the analysis of abstract data I decided to work with smaller individual items on a larger scale so as to keep the analysis itself simpler to implement and run. This led me to working with Tweets.

The learning curve as I embarked on the project was far from lenient. At

several points considerably steeper than previous work I had undertaken. My previous knowledge of Python did not cater for the scale of this work and my knowledge of raw text analysis was, at the time non-existent. Throughout development I have learnt everything necessary about these subjects and how they can be applied to accurate Text Analysis.

## **1.2 Goals**

The primary goal of the project was to build a program capable of simultaneously running efficient and effective data collection and analysis while also capable of compiling all previously assessed data for analysis and aggregation on demand. This data was to be highly generalized, easily readable for both human users and predictive analysis.

# Chapter 2

## Analysis

### 2.1 Objectives

The main objective of the project is to create an effective and efficient text analysis, capable of being performed in real time for any variety of uses, from market research to predictive analysis. The analysis was kept generalized as topical language development would have taken a much greater length of time. For example if searching for opinions on a horror movie, words such as terrifying and gory would be positive, however in a generalised case these would be counted as negative opinions. As such the idea of topical text analysis was dropped before development even began.

Scalability was a topic discussed often while I was working at IBM. To ensure ease of scaling and parallelization I developed the project exclusively on my own laptop. As such some later basic shortcomings and limits in processing speed could be compensated in a higher level implementation. However the basic implementation on the laptop had to function effectively as a proof of concept.



# Chapter 3

## Design

The final version of the project was composed of 6 main modules to run. The main command line interface WOPMain.py[4], the text cleaner textClean.py[4], the text analyzer textAnalysis.py[4], the tweet handler tweetHandler.py[4], the tweet scraper twitterScraper.py[4] and a database class MyDB.py[4].

### 3.1 Main Interface

Once the main interface is started it gets the details of all current clients from the database. Within the database clients are stored as strings used to scrape data, coupled with a rate in seconds between data collection and a count of tweets per collection. For each client a thread is created that waits for the specified seconds before attempting to collect data from Twitter. An entry is also made into the log on the initialisation of thread as well as the current wait time. All threads share a single database interface and will attempt to complete their wait between Twitter calls if the main is shut down. Each thread independently requests, analyzes and stores tweets until the main is closed.

### 3.2 Text Cleaner

The text cleaner was designed to act as an early level filter for the project. It was designed to clear away links as well as irrelevant tweets. Tweets wherein links made up more than half of the text were dropped before analysis took place. Initially this was supposed to make use of Twitter entities, however due to an issue in gathering older testing data and the projects workaround, this functionality was abandoned.

### **3.3 Twitter Scraper**

The scraper handles all calls to the official Twitter API. Initially it was built using the TwitterSearch[5] library, however this library led to several issues late in development when threading was being implemented. To resolve this problem development was moved to the Twython[6] Library. Both of these libraries were officially licensed by Twitter. Twitters API consistently caused issues throughout development due to a number of factors which will be covered in Evaluation, see Page 5.

### **3.4 Twitter Handler**

The handler was built to parse the key information from Twitter API calls. It was initially created early on in development however it was then dropped from the project and rebuilt towards the end of development. Each individual tweet was comprised of roughly 200 lines of different JSON values, some of which were depreciated so parsing the key details and reducing raw data flow was a key part of the overall process.

### **3.5 Text Analyzer**

The text analyzer was one of the two main focuses of the project. It was specially designed for Twitter, using small features such as emojis and hash-tags to build a more accurate, in-depth analysis. One of the most important aspects of this process, however was to streamline it as a real world implementation would likely deal with hundreds of thousands of tweets per minute. At the time of this report the average analysis time for the text of a tweets took less than a ten thousandth of a second.

### **3.6 Database**

The database class was built using Object Oriented Python to encapsulate all Database functionality for the project. Initially I had planned on developing the project around a MongoDB database due to the JSON dictionary structure of tweet information. However as development moved ahead the data became more reduced and concise, lending itself more to a MySQL implementation than that of MongoDB. The class was used for all database interaction, from simple data request to aggregating tweet information for graphing.

# Chapter 4

## Implementation

### 4.1 Project Planning




Figure 4.1: First Gantt Chart Draft

Project Planning began immediately after the project brief was accepted in October. As can be seen above 4.1 there an initial estimate of 6 weeks necessary for research and planning. With work to begin with the Twitter API towards the end of the sixth week. A five week period from the end of November through December was set aside to allow time for the Christmas Exams as well as the Christmas break. Work on interfacing with the API was to resume immediately after New Years and work on the text analysis was to begin around the same time. Testing on various small project elements was to take place throughout development from assessing Python Libraries for Twitter to resting more efficient means of Text Analysis.

As is often the case in any planned project, the initial plan was the first thing to go. Early Twitter API test went so well I began work on a Python interface a week earlier than planned. After Christmas I managed to finish up work on my TwitterSearch[5] wrapper module and began work on the Text Analysis. Before Christmas I had done research into different Text Analysis Software such as Natural Language Toolkit or NLTK?? and began with a simple unweighted words scan. Comparing words to items in predetermined




Figure 4.2: Final Gantt Chart Draft

.txt files containing positive and negative words. It was slow and clunky, but functional. Satisfied for now I began planning the Database Structure.

## 4.2 Disecting Tweet Structure

Early work in development began by dumping collected tweets in a JSON[8] file. Each tweet consisted of nearly 200 different tags and values. As a result parsing the usable information from the structure proved to be difficult as large amounts of data were repeated. Furthermore several of the tags used were depreciated and served no purpose whatsoever.

## 4.3 Database Structuting

Due to the JSON[8] structure of the tweets I initially began implementing a MongoDB database on a topic by topic basis. However as I began giving more thought to the overall database structure and how best to minimize cost while maximizing usable data and reducing data aggregation complexity I realised

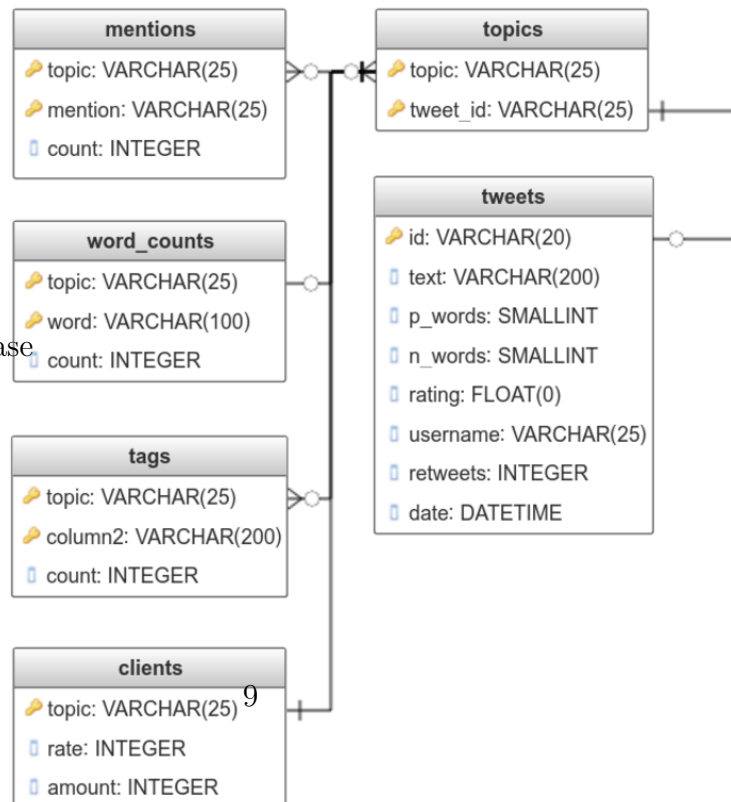


Figure 4.3: Finalized Database Schema

that I had simplified by data to the extent that it could be reduced to a more cost effective MySQL database.

## 4.4 Test Data Collection

Following the database restructuring I continued work on the Text Analysis, however I found myself lacking in ideas for usable topics and data. In the end I decided to build a several large collections of Tweets from the previous 3 U.S. Presidential Elections. I ran a request through the official API but discovered that it limited returned Tweets to those less than a week old. I began research and found an unofficial library that scraped data directly from the Twitter Advanced Search Browser API. I began with a quick collection of ten thousand tweets from the thirty days prior to the actual election day for each candidate. (Barack Obama twice). At the suggestion of a fellow student I also pulled tweets on Sarah Palin for the 2008 U.S. Presidential Election as she was just as large a point of discussion as John McCain. This process ultimately left me with about sixty thousand tweets. However, disappointed in the variety of data I wrote a small script to pull up to ten thousand tweets from each day up to thirty one days prior to the actual election.

This process ultimately led to several complications. The largest of which was a combination of Twitters API and the unofficial Python library. As I was making multiple requests to the API for large volumes of data, if I made too many my access to the API would be suspended by Twitter. Secondly the library was designed such that if a request to the API was rejected it would exit Python, breaking any loop in which I attempted use for the process. Ultimately I changed the library myself to remove this feature and successfully created a functional script. The script left an hour wait in between API calls as well as in the event of an API rejection. After a little under two weeks and minimal API rejections I finally had a satisfactory testing dataset.

## 4.5 Text Analysis

While I ran the data collection script I began enhancing the Text Analysis, allowing for functional words to invert, enhance or detract from the weight a word had on the overall rating of a tweet. These function words also affected each other but the weight would be ignored if a positive or negative word did not follow them. Hashtags were also added as a means of emphasis on whatever word they were used and mentions were catalogued to allow for analysis of individual users. As a final addition a small library of emojis were added to the analysis. Ultimately the greatest enhancement to the text analysis came in the form of a small fix. Throughout most of development words were stored in a list with a linear search through each list used to inform the program of a words classification. However simply replacing this feature with a simple try/catch and dictionaries increased the speed of text analysis ten-fold.

Following most of my work on the Text Analysis I began working on a means of interacting with and controlling the software. I began with a small and basic interface that ran each client in a queue. However this implementation was incredibly inefficient so I began working on running each client on individual threads, something I had never dealt with in Python. This took a little over a week to finalize and hit multiple setbacks during implementation. One of the first issues was in safely shutting down the interface. Due to the delay in client request, (rate in the database table) the interface itself could not shut down until the delay had completed for all threads. A larger problem however, was an issue surrounding the `TwitterSearch[5]` Python library, which had been causing a series of errors with the threads, ultimately preventing any threads from successfully starting. To resolve this issue I abandoned the `TwitterSearch[5]` library for the `Twython[6]` library which thankfully ran smoothly when added to the threaded implementation.

Most of the later development on the Interface was carried out in tandem with later database development. As multiple UI features were added to simplify the interface for a generalized user as well as several commands for data aggregation and display. Some of this more complicated data compilation did prove to be the larger of issues within the final stages of development

# Chapter 5

## Evaluation

Overall, I am unsatisfied with how my project ended up. There are many elements I would have loved to have gotten the opportunity to refine further which I may ultimately end up doing in my own spare time.

### 5.1 Weights

At present there is a bare minimum weight function built into the interface that weighs the overall rating of a tweet based on its retweets in comparison to the daily or hourly average, depending on data grouping.




Figure 5.1: Graphed Data

Unfortunately due to the low volume retweets on the majority of database entries the resultant value becomes  $\approx 0$ , resulting in the data being ignored. As shown in Figure 5.1(a) as opposed to Figure 5.1(b). I attempted to get some data to show the comparison at a daily scale, however when weighted there was no available data as opposed to the unweighted data in Figure 5.1(c).

In the end the function removed most available data for any given topic with most tweets falling far below this average. Given more time and further development it would have been interesting to look into a more accurate and in-depth means of weighing these tweets for data compilation. However at present a concrete implementation for doing so escapes me.

## 5.2 Data Collection

As a proof of concept I decided early on Throughout development Twitter itself proved to be the largest bottleneck. The official API only allows access to data less than a week old and limits individual request to 100 individual tweets. On top of this there are additional limits on how many request can be made from a developer and/or IP address[9]. One of the largest delays in the projects overall development was the time spent collecting data from the 2008, 2012 and 2016 elections as I had to use an unofficial library to scrape the data. This was both inefficient and was contrary to one of my open source development restriction. However I ignored this as the library itself used an MIT License[2] and was not a part of the final project. Furthermore I had to alter the code of this library myself to allow for pausing my script in the event of an API rejection. The data I received from this library was also inefficient as it lacked many of the more specialised tags included in officially retrieved twitter data, such as the entities tags which would have allowed for faster retrieval of mentions, hashtags and most importantly links. Due to this limitation I had to reduce the functionality of the text cleaning module. I had initially planned on having the module access the link and replace the link with the title of the web page within the text body of the tweet.

Looking back over the project now, were I to do the project all over again I definitely would have looked into more Twitter libraries, or created my own means of scraping the data so that I would be far less limited by the official Twitter APIs.



# Chapter 6

## Conclusions

### 6.1 Development

Overall I feel the development process was ultimately held back by the overall

# Bibliography

- [1] Python 2.7 Documentation <https://docs.python.org/2/>
- [2] MIT Open Source License <https://opensource.org/licences/MIT>
- [3] IBM Irish Website <https://www.ibm.com/ie-en>
- [4] Project Source Code GITHUB LINK TBA
- [5] TwitterSearch Repository <https://github.com/ckoepp/TwitterSearch>
- [6] Twython Repository <https://github.com/ryanmcgrath/twython>
- [7] Natural Language Toolkit [www.nltk.org](http://www.nltk.org)
- [8] JSON [www.json.org](http://www.json.org)
- [9] Twitter API Limits <https://dev.twitter.com/rest/public/rate-limits>